

Database Management Systems

Lorraine Goeuriot

IUT 1, Université Grenoble Alpes

Auteur du cours : Marie-Christine Fauvet

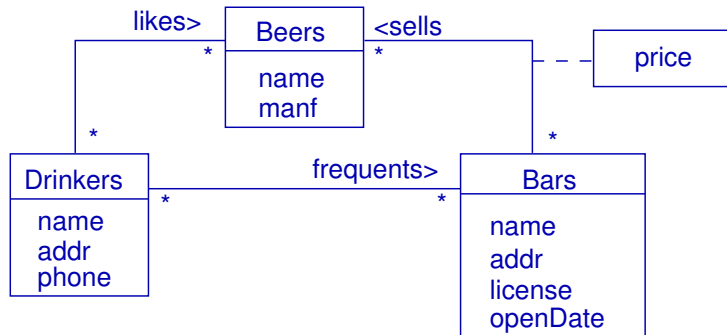
Université Joseph Fourier, Grenoble – UFR IM²AG

21 janvier 2016

Contenu

- 1 SQL - Data Analysis
 - Introduction
 - SQL Syntax
 - Clause Select
 - Grouping relations
 - Set operations
 - Embedded queries
 - Agrégation et partition
 - Embedded queries - advanced
 - Empty values

New Toy DB¹



1. This DB is taken from : *Database Systems : The Complete Book (2e édition)*, écrit par Hector Garcia-Molina, Jeff Ullman, and Jennifer Widom, 2008.

Relations

Drinkers (*name*, *addr*, *phone*)

Beers (*name*, *manf*)

Bars (*name*, *addr*, *license*, *openDate*)

Likes (*drinker*, *beer*)

Likes[*drinker*] \subseteq *Drinkers*[*name*]

Likes[*beer*] \subseteq *Beers*[*name*]

Sells (*bar*, *beer*, *price*)

price > 0

Sells[*beer*] \subseteq *Beers*[*name*]

Sells[*bar*] \subseteq *Bars*[*name*]

Frequents (*drinker*, *bar*)

Frequents[*drinker*] \subseteq *Drinkers*[*name*]

Frequents[*bar*] \subseteq *Bars*[*name*]

Domains should be described too...

Relations Values

Bars

Name	Addr	License	openDate
Australia Hotel	The Rocks	123456	12/1/1940
Coogee Bay Hotel	Coogee	966500	31/8/1980
Lord Nelson	The Rocks	123888	11/11/1920
Marble Bar	Sydney	122123	1/4/2001
Regent Hotel	Kingsford	987654	29/2/2000
Rose Bay Hotel	Rose Bay	966410	31/8/2000
Royal Hotel	Randwick	938500	26/6/1986

Relations Values

Drinkers

Name	Addr	Phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321
Marie	Rose Bay	9371-2126
Adrian	Redfern	9371-1244

Frequents

Drinker	Bar
Adam	Coogee Bay Hotel
Gernot	Lord Nelson
John	Coogee Bay Hotel
John	Lord Nelson
John	Australia Hotel
Justin	Regent Hotel
Justin	Marble Bar
Marie	Rose Bay Hotel

Relations Values

Beers

Name	Manf
80/-	Caledonian
Bigfoot Barley	Sierra Nevada
- Wine	
Burraborang	George IV
- Bock	- Inn
Crown Lager	Carlton
Fosters Lager	Carlton
Invalid Stout	Carlton
Melbourne Bitter	Carlton
New	Toohey's
Old T	oohey's
Old Admiral	Lord Nelson
Pale Ale	Sierra Nevada

Premium Lager	Cascade
Red	Toohey's
Sheaf Stout	Toohey's
Sparkling Ale	Cooper's
Stout	Cooper's
Three Sheets	Lord Nelson
Victoria Bitter	Carlton

Relations Values

Likes	
Drinker	Beer
Adam	Crown Lager
Adam	Fosters
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

Relations Values

Sells		
Bar	Beer	Price
Australia Hotel	Burraborang Bock	3.50
Coogee Bay Hotel	New	2.30
Coogee Bay Hotel	Old	2.50
Coogee Bay Hotel	Sparkling Ale	2.80
Coogee Bay Hotel	Victoria Bitter	2.30

Query - Definition

A query is a *declarative program* aiming at retrieving data from a DB.
There are two ways to use them :

- Interpreted : directly query the database (e.g. via terminal or DBMS).
- Encapsulated : the query and results are treated by a programme (e.g. web application).

Example :

```
select bar, price /* a set of attributes */  
from Sells /* a set of relations */  
where beer = 'Victoria Bitter' /* a boolean expression */
```

*In which bar and at which price the beer
'Victoria Bitter' is sold.*

With MySQL...

```
SQL> select bar, price  
2 from Sells  
3 where beer = 'Victoria Bitter';
```

<i>BAR</i>	<i>PRICE</i>
<i>Coogee Bay Hotel</i>	<i>2.3</i>
<i>Marble Bar</i>	<i>2.8</i>
<i>Regent Hotel</i>	<i>2.2</i>
<i>Royal Hotel</i>	<i>2.3</i>

```
4 rows selected  
SQL>
```

SQL Identifiers

- Objects : relations, attributes are identified by their names, with conventions similar to usual programming languages, BUT names are not case sensitive.
Examples : Drinkers, Beer, select, From, wherE, And, OR, etc...
- Constants : string and number values have an implicit type.
Examples : 'C\' is a string, " John's bike", 123.84, 4.23e-12, .01, 23., etc...

Types de données SQL

- Strings :

char [n] : n characters (fixed length)

varchar [n] : 0..n char (variable length)

- Numbers :

smallint, integer, bigint

integer (p)

real, float, double

float(p)

- Sets and tuples : list of values between (and)

Examples : (12, 3, 9, 10) is a set (could be a tuple!)

(123, 'John', '12-3-2001') is a tuple

http://www.w3schools.com/sql/sql_datatypes_general.asp

- Date, Time, Datetime... : various format for temporal types.

Example :

- *'2015-01-13' type Date*
- *'15 :30 :01' type Time*
- *'2015-01-13 15 :30 :01' type Datetime*

Operators and functions

- Comparison operators (work with every type) :
<, >, <>, !=, =
- Boolean operators :
NOT, AND, OR (from highest to lowest priority)
- Arithmetic operators :
+ sum, - subtraction, * multiplication, / division
- Other numerical operators :
abs, ceil, floor, truncate, round, exp, pow, mod, sqrt, etc...
sin, cos, tan, atan, etc...

- String functions :

concat, lower, upper, substring, length, etc...

- Temporal functions :

to_date (string,string) → date

/ to_date ('05 Dec 2001', 'DD Mon YYYY') */*

to_char (date, string) → string

/ to_char (d,f) gives a string representing d following format f. */*

/ to_char (to_date ('05 Dec 2001', 'DD Mon YYYY'), 'DD/MM/YY') = '05/12/01' */*

Other temporal functions :

add_months (date, integer) → date

*/*add_months (d1, n) gives date d1 plus one month.*/*

Numerical operators can be used on temporal values :

d1, d2 : date, n : int ≥ 0

d1 - d2 amount of days between d1 and d2

d1 + n is the date n days after d1

d1 - n is the date n days before d1

Other functions :

sysdate → date, etc...

Projection, selection

Give the bars (name and address) located in Sydney

In SQL :

```
select name, addr /*projection*/  
from Bars /*relation(s)*/  
where addr = 'Sydney' /*selection*/
```

In the query *select ... from ... where P*, P is a predicate, built with :

- a simple condition, such as :
 <att. name> <comp. op.> <att. name> ou
 <att. name> <comp. op.> <cst. val.>
- or a complex boolean expression such as :
 <cond.> <bool. op.> <cond.>
 where <cond.> is a simple condition or a complex boolean expression.

Comparison operators : =, \neq , <, >, etc.

Boolean operators : not, and, or

P can contain brackets

Exercises

1) Evaluate the following expressions :

- A and B or C
- A and (B or C)
- not A and A

when A=false, B=true, C=false

2) Evaluate the following expressions : :

- $(3 > 5)$ and $(6 = 6)$
- $(3 > 5)$ or $(6 = 6)$
- $(X \geq Y)$ and $(B \text{ or } (X = 3))$, when $X=4$ et $Y=3$

Expressions in a Select clause

For each bar, give its name, address and for how many years it has been open

```
select name, addr,  
       to_number (to_char (sysdate, 'YYYY'))  
       - to_number (to_char (openDate, 'YYYY'))  
from Bars;
```

Many operators and functions exist :

<https://dev.mysql.com/doc/refman/5.7/en/functions.html>

Multiple Values

Give the beers that cost less than \$2.5

The expected result is :

Beer
New
Old
Victoria Bitter

In SQL :

*select Beer from Sells
where price <= 2.5*

And the result is :

Beer
New
Old
Victoria Bitter
New
Victoria Bitter
New
Old
Victoria Bitter

Multiple Values - get rid of them

To remove repetitions from the results :

***select distinct** Beer from Sells where price \leq 2.5*

Use it with caution :

- Think about the cost !
- Often useless : *select name from Drinkers*

Use of the star (*)

The character * means “no projection”

Example : Give all the information about the drinkers In SQL : *select * from Drinkers*

Result :

name	addr	phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321
Marie	Rose Bay	9371-2126
Adrian	Redfern	9371-1244

- Very useful to visualize the schema and content of a table

Rename attributes in a *select*

Attributes can be renamed in a query using *AS* in the *select* clause :

select name as Drinkers, addr, phone from Drinkers

Result :

drinkers	addr	phone
Adam	Randwick	9385-4444
Gernot	Newtown	9415-3378
John	Clovelly	9665-1234
Justin	Mosman	9845-4321
Marie	Rose Bay	9371-2126
Adrian	Redfern	9371-1244

Order tuples

The *order by* clause can only be applied to attributes in the *select* clause :

select bar, beer, price from Sells

order by price desc, bar asc

Result :

bar	beer	price
Lord Nelson	Three Sheets	3.75
Lord Nelson	Old Admiral	3.75
Australia Hotel	Burraborang Bock	3.5
Coogee Bay Hotel	Sparkling Ale	2.8
Marble Bar	New	2.8
Marble Bar	Victoria Bitter	2.8
Marble Bar	Old	2.8
Coogee Bay Hotel	Old	2.5
...		

Relational junction (intern)

```
select A1, A2, ..., An           /* → projection*/  
from R1 join R2 on .. join Rp on ... /* → relation grouping*/  
where C                          /* → selection*/
```

For each beer sold by Australia Hotel, give its price and manufacturer

In SQL :

```
select Beer, Price, Manf         /* projection */  
from Beers join Sells on (Name = Beer) /* join condition */  
where Bar = 'Australia Hotel'    /* selection condition */  
/* The clause "on" can contain any predicate. */
```

To join two (or more) relations

Name ambiguity in a query

How can we handle several occurrences of the same name in a query ?

For each bar that John frequents, give its name, the beers it sells and their price.

```
select Bar, Beer, Price  
from Sells join Frequents (Bar = Bar)  
where Drinker = 'John'  
ERROR at line 2 :  
ORA-00918 : column ambiguously defined
```

Attribute name ambiguity - solution

The complete name of an attribute is characterized by the relation name :

```
select Sells.Bar, Sells.Beer, Sells.Price  
from Sells join Frequents (Sells.Bar = Frequents.Bar)  
where Frequents.Drinker = 'John'
```

- Each attribute is characterized by the relation it belongs to (easier to read)
- Necessary for the attributes defined in more than one relation in the from clause.

Natural join : first form

The join condition relates to all the common attributes of the 2 relations.

```
select Bar, Sells.Beer, Sells.Price  
/* Bar appears only once in the resulting schema */  
from Sells natural join Frequent  
/* The attribute bar is common to both relations */  
where Frequent.Drinker = 'John'
```

The Bar attribute (appearing in both Sells and Frequent) cannot be prefixed by the relation name.

Natural join : 2nd form

The condition does not relate to all the common attributes

```
select Bars.name, Addr, Drinkers.name
from Bars join Drinkers using (addr)
/* the join condition uses only the common attribute addr*/
```

Another version :

```
select Bars.name, Bars.Addr, Drinkers.name
from Bars join Drinkers on (Bars.addr=Drinkers.addr)
/* The join condition is explicit.*/
```

Natural join and relational join in the same clause from

For each beer, give its name, manufacturer, the name of drinkers who like it, and for each of these drinkers the bars they frequent.

```
select name, manf, drinker, bar
```

```
from Beers join Likes on (name=beer) natural join Frequents
```

/* The following expression is not correct : the drinker attribute is ambiguous : it belongs to both relations Frequents natural join Beers and Likes.

```
select name, manf, drinker, bar
```

```
from Frequents natural join Beers join Likes on (name=beer)
```

/* The following expression is correct : */

```
select name, manf, drinker, bar
```

```
from Frequents natural join (Beers join Likes on (name=beer))
```

Join operations all have the same priority, they are evaluated from left to right.

Parenthesis allow to force the order of evaluations.

Cross join

```
select ....  
from Bars cross join Drinkers  
where Drinkers.name = 'John'
```

Join : summary

R and S are relations defined as : R (X, Y, Z) et S (Y, Z, T)

- Cross join : **from R cross join S**
schema : **R.X, R.Y, R.Z, S.Y, S.Z, S.T**
- Relational join : **from R join S on (P)**
schema : **R.X, R.Y, R.Z, S.Y, S.Z, S.T**
P is a predicate, that must be True on **R.X, R.Y, R.Z, S.Y, S.Z, S.T**
The join condition is **P**
- Natural join (1st form) : **from R natural join S**
schema : **R.X, Y, Z, S.T**
The join condition is **R.Y = S.Y and R.Z = S.Z** (All the common attributes are concerned)
- Natural join (2nd form) : **from R join S using (Y)**
schema : **R.X, Y, R.Z, S.Z, S.T**
The join condition is **R.Y = S.Y**

Multiple use of a relation

Give all the pairs drinkers who like the same beer

```
select L1.drinker, L2.Drinker
from Likes L1 join Likes L2
    on (L1.beer = L2.beer and L1.drinker <> L2.drinker)
```

Questions :

Is the clause `select distinct` necessary to remove duplicates ?

How can we remove symmetrical pairs?

R is antisymmetrical if :

$\langle X, Y \rangle \in R \implies (\langle Y, X \rangle \notin R \text{ or } X = Y)$

We know that for all X, $\langle X, X \rangle \notin R$

```
select distinct L1.drinker, L2.Drinker
from Likes L1 join Likes L2
on (L1.Beer = L2.Beer and L1.Drinker < L2.Drinker)
```

Union, Intersection, Difference

Q1 and Q2 are queries with a similar structure `select... from....` with compatible schema :

- `select A, B from...` is not compatible with `select C from...`
- `select A, B from...` is compatible with `select C, D from...` if and only if A and C are comparables, as well as B and D.

Opérateurs

- Union [all]
Duplicates are not removed with option all
- Intersect (intersection)
- Minus (différence)

Duplicates are removed, except with union all.

Difference : example

Give the beers sold for less than \$3, that John does not like.

select Beer from Sells where Price < 3

minus

select Beer from Likes where Drinker = 'John'

Both subquery must create relations with compatibles schema.

Intersection : example

Give the drinkers and the beers such as drinkers like a given beer and frequent a bar that sells it.

```
select Drinker, Beer from Likes  
intersect  
select Drinker, Beer from Sells natural join Frequents
```

Another version :

```
select distinct Drinker, Beer  
from Likes natural join Sells natural join Frequents  
/* Joins are evaluated from left to right. */
```


Union : example

Give the drinkers who like Sparkling Ale or frequent Lord Nelson.

```
select Drinker from Likes  
where Beer = 'Sparkling Ale'  
union  
select Drinker from Frequents  
where Bar = 'Lord Nelson' →
```

Drinker
Gernot
Justin
John

```
select Drinker from Likes  
where Beer = 'Sparkling Ale'  
union all  
select Drinker from Frequents  
where Bar = 'Lord Nelson' →
```

Drinker
Gernot
Gernot
Justin
John

Queries in the clause from

The clause `from` can contain SQL expressions.

Give the name and manufacturer of the beer sold less than \$3 that John does not like

We already saw the query giving the beers sold less than \$3 that John does not like :

```
select Beer from Sells where Price < 3  
minus
```

```
select Beer from Likes where Drinker = 'John'
```

We call `R` the relation resulting from this query.

The final query is (using R) :

```
select Beer, Manf  
from Beers join R on (Beer = Name)
```

If we substitute R to its expression :

```
select Beer, Manf  
from Beers join (select Beer from Sells where Price < 3  
                 minus  
                 select Beer from Likes where Drinker = 'John') R  
  /* < b > ∈ R ⇔ b is sold less than $3  
   et John n'aime pas b. */  
on (Beer = Name)
```

A name has to be given to the subquery, also called nested query (R in this case).

Why are embedded queries useful ?

- Decompose the query in simpler subqueries.
- The expression and test of the subqueries are independent.

Rule : the relation embedded in the clause from MUST be specified.

Clumsy use :

```
select beer, price
from (select distinct beer from Sells where price  $\geq$  10) X
/*  $\langle b \rangle \in X \iff$  there is a bar selling beer b for more than $10. */
natural join Sells
where bar = 'Coogee Bay Hotel'
```

A shorter and simpler version :

```
select beer, price from Sells
where price  $\geq$  10 and bar = 'Coogee Bay Hotel'
```

Aggregate Functions

Used to reduce a list of values to a single value.

- Count (*) → number of tuples
- Count (A) → number of values for A
- Count (distinct A) → number of distinct values for A
- Avg (A) → average value for A
- Min (A) (resp. Max) → minimum (resp. maximum) value of A
- Sum (A) → sum of values for A

Aggregate functions : Example

What is the average price for beers sold in Australia Hotel ?

select avg (price) from Sells where Bar = 'Australia Hotel'

How many bars are located in The Rocks ?

select count () from Bars where Addr = 'The Rocks'*

Aggregate functions : more examples

How many bars are selling beers for less than \$2.5 ?

select count (Bar) from Sells where price < 2.5

Careful : This query is not correct, why ?

A bar is counted as many times as it is selling beers for less than \$2.5

Correct query :

select count (distinct Bar) from Sells where price < 2.5

Which bars sell beer New for the lowest price ?

select bar

from Sells join

(select min(price) as minP from Sells

where beer='New') Min

/ <m> ∈ Min ⇔ m is the lowest price for beer New
among all the bars.*/*

on (price = minP)

where beer='New'

Partition

Partition a query to apply an aggregate function on each class separately
How many beers do drinkers like ?

Expected result :

Drinker	Beer
Adam	3
Gernot	2
John	4
Justin	2

Partition - creation

1. Create a partition on Likes

Adam	Crown Lager
Adam	Fosters
Adam	New
Gernot	Premium Lager
Gernot	Sparkling Ale
John	80/-
John	Bigfoot Barley Wine
John	Pale Ale
John	Three Sheets
Justin	Sparkling Ale
Justin	Victoria Bitter

In SQL :

```
select ...  
from  
Likes  
group by  
Drinker
```

2. *Reduce each class of the partition to a single tuple*

Drinker list (partition criteria)

→ *Drinker (one of the values)*

Beer list

→ *int (# values)*

In SQL :

select Drinker, count (Beer)

from Likes

group by Drinker

Proper use of the clause *select*

In a query containing the clause *group by*, *select* only contains :

- One or several attributes among the partition criteria
- One or several aggregations applied to the attributes

Wrong :

```
select Drinker, Addr, count (Beer)
from Likes join Drinkers
on (Drinker=Name)
group by Drinker
```

Right :

```
select Drinker, Addr, count (Beer)
from Likes join Drinkers
on (Drinker=Name)
group by Drinker, Addr
```

Partition filter

For each bar selling more than 2 beers, give the number of beers it is selling and the number of drinkers frequenting it.

```
select Bar,  
       count (distinct S.Beer) as NbBeers,  
       count (distinct F.Drinker) as NbDrinkers  
from Sells S natural join Frequents F  
group by Bar  
having count (distinct S.Beer) > 2
```

More examples...

Which bars are selling all the beers ?

We know that : $|A| = |B| \wedge A \subseteq B \implies A = B$

For each bar, how many beers ?

```
select bar, count(beer) as nbBeers from Sells group by bar
```

How many beers sold in total ?

```
select count(distinct beer) as nbTot from Sells
```

The final query is :

```
select bar
from (select bar, count(beer) as nbBeers from Sells group by bar) X1
join (select count(distinct beer) as nbTot from Sells) X2 on
(nbBeers = nbTot)
```

Which bars are selling the maximum number of beers ?

- For each bar, how many beers ? see X1 above.
- In X1, which bar is associated to the maximum ?

```
select bar from X1 join (select max(nbBeers) as nbMax from X1) Y
on (nbBeers = nbMax)
```

The final query :

```
select bar
from (select bar, count(beer) as nbBeers from Sells group by bar) X1
join (select max(nbBeers) as nbMax from X1) Y
on (nbBeers = nbMax)
```

For each drinker who likes all the beers sold in Australia Hotel, give the name and the bars she/he likes.

```
/* How many beers are sold in Australia hotel ? */
```

```
select count(Beer) from Sells where Bar = 'Australia Hotel'
```

```
/* How many beers each drinker likes ? */
```

```
select drinker, count(beer) as nbBeers from Likes group by Drinker
```

```
/* Final query : */
```

```
select drinker, bar
```

```
from (select count(Beer) as nbBeers
```

```
      from Sells where Bar = 'Australia Hotel') X1
```

```
      natural join
```

```
      (select drinker, count(beer) as nbBeers
```

```
      from Likes) X2
```

```
      natural join Frequents
```

This query is not correct, why ?

Correct query :

```
select drinker, bar
from (select count(Beer) as nbBeers
      from Sells where bar = 'Australia Hotel') X1
natural join
(select drinker, count(beer) as nbBeers
 from Likes natural join Sells
 where bar = 'Australia Hotel') X2
natural join Frequents
```


Evaluation order in queries

select 5. *projections or aggregate functions*

select 5. projection on a partition criteria and aggregate functions

from 1. *relation junction*

where 2. *selection on the joined relations*

group by 3. partition

having 4. filter on the partition

Partition : common error

For each bar, give its name, address, the amount of beers it sells, and the drinkers frequenting it.

```
select bar, count (distinct beer) as nbBeers, addr, drinker
*
from Bars join Sells on (name = bar)
        natural join Frequents
group by bar
```

ERROR at line 1 :

ORA-00979 : not a GROUP BY expression

Why ?

First problem : attribute Addr

```
select bar, count (beer) as nbBeers, addr
/* distinct is not necessary. */
from Bars join Sells on (name = bar)
group by bar, addr
```

Expressions *group by name* and *group by name, addr* relate to the same partition, as a bar has only one address.

Result

Bar	nbBeers	Addr	Drinker
Australia Hotel	1	The Rocks	{John}
Coogee Bay Hotel	4	Coogee	{Adam, John}
Lord Nelson	2	The Rocks	{Gernot, John}
Marble Bar	3	Sydney	{Justin}
Regent Hotel	2	Kingsford	{Justin}

Column Drinker's type is a set !

Result

Bar	nbBeers	Addr	Drinker
Australia Hotel	1	The Rocks	John
Coogee Bay Hotel	4	Coogee	Adam
Coogee Bay Hotel	4	Coogee	John
Lord Nelson	2	The Rocks	Gernot
Lord Nelson	2	The Rocks	John
Marble Bar	3	Sydney	Justin
Regent Hotel	2	Kingsford	Justin

Pour résoudre le problème sur Drinker

```
select bar, nbBeers, addr, drinker
from (select bar, count (beer) as nbBeers, addr
      from Bars join Sells on (name = bar)
      group by bar, addr) X
/* <b, n, a> ∈ X  $\iff$  bar b sells n beers and its
address is a.*/
natural join Frequents
```

Operator IN

- *The result of a query is a relation.*
- tuple IN relation \iff tuple \in relation
- Inverse : tuple NOT IN relation

Give the name and manufacturer for beers that John likes.

```
select Name, Manf from Beers
where Name in (select Beer from Likes
               where Drinker = 'John')
```

Another expression :

```
select B.name, B.manf
from Beers B join Likes L on (B.name = L.beer)
where L.drinker = 'John'
```

Operator IN - syntax

The operator IN can be generalized to deal with tuples of any length

- $\langle X_1, X_2, \dots, X_n \rangle \text{ IN relation (with attributes } A_1, A_2, \dots, A_n) \text{ is true}$
 $\iff \langle X_1, X_2, \dots, X_n \rangle \in \text{relation.}$

select ... from ...

where (X_1, X_2, \dots, X_n) in

(select A_1, A_2, \dots, A_n from)

*Obviously, for each $i=1..n$,
types for X_i et A_i must be comparable*

Use of nested queries

Give the bars that serve New at the same price as Victoria Bitter in Coogee Bay Hotel.

```
select Bar from Sells
where Beer = 'New' and Price =
    (select Price from Sells
     where Bar = 'Coogee Bay Hotel' and Beer = 'Victoria Bitter')
```

If the nested query gives a single tuple, the system gives it back, else it returns an error

Mieux :

```
select Bar from Sells
where Beer = 'New' and Price IN
    (select Price from Sells
     where Bar = 'Coogee Bay Hotel' and Beer = 'Victoria Bitter')
```

Utilisation de sous-requêtes

```
select Bar from Sells
where Beer = 'New' and Price =
    (select Price from Sells
     where Bar = 'Coogee Bay Hotel' and Beer = 'Victoria Bitter')
```

Better :

```
select Bar from Sells
where Beer = 'New' and Price IN
    (select Price from Sells
     where Bar = 'Coogee Bay Hotel' and Beer = 'Victoria Bitter')
```

Even better :

```
select s1.bar from Sells s1 join Sells s2 on (s1.price=s2.price)
where s1.beer = 'New' and
    s2.bar = 'Coogee Bay Hotel' and s2.beer = 'Victoria Bitter'
```

Parameters range

Range rule : Parameters can be used in any nested query
For each bar, give the cheapest beer and its price (in this bar)

```
select Bar, Beer, Price
from Sells S1 /* S1 is a parameter */
where Price in (select min (Price)
               from Sells S2 where S1.Bar = S2.Bar)
```

Another expression :

```
select Bar, Beer, Price
from Sells S join
  (select bar, min(price) as minPrice
   from Sells
   group by bar) as X
on (S.bar=X.bar and price=minPrice)
```

Motivation

There can be a need to use empty values :

- Temporarily unknown values :
I don't know this address yet, but will know it soon
- Temporarily forbidden values :
She's not married, she can't have a marital name yet
- Strictly forbidden values :
He didn't pass the exam, so can't have a mark.
- etc...

From DBs perspective, there is only one "empty value" concept.

Empty value characteristics

- The empty value is called NULL.
- NULL values don't appear in the queries output (nothing appears).
- NULL belongs to all types.
- All the operators have a specific behaviour with empty values.
- Two new operators : IS NULL - IS NOT NULL

Operators and empty values

- Comparison operators : A is an attribute
 $\forall \text{ op} \in \{<, >, <=>, =\}$, A op Null, Null op A equals Null
- Arithmetic operators :
 $\forall \text{ op} \in \{+, -, /, *, \dots\}$, A op Null, Null op A equals Null
- Boolean operators :

AND	True	False	Null
True	True	False	Null
False	False	False	False
Null	Null	False	Null

OR	True	False	Null
True	True	True	True
False	True	False	Null
Null	True	Null	Null

Similarly, NOT null equals null

- min, max, sum, avg ignore NULL values.
If all the values are NULL, the result is NULL.
- count(*) doesn't take into account NULL values : it only counts tuples.
- count (distinct A) et count (A) only count non NULL values.

*NULL isn't 0, { }, "", nor ()
NULL is nothing!*

Queries and missing values

R	Stuld	Mark	Course
	12	45	maths
	11	90	french
	11	75	maths
	12	65	physics
	12		french

Give the information on the students who got a mark different from 90

select * from R where note <> 90

R	Stuld	Mark	Course
	12	45	maths
	11	75	maths
	12	65	physics

*A tuple is selected if it satisfies the where condition
(if the evaluated condition is True)*

Give the information on the students who got a mark different from 90 (including those who didn't pass the exam)

select * from R where note <> 90 or note is null

R	Stuld	Mark	Course
	12	45	maths
	11	75	maths
	12	65	physics
	12		french

For each student, calculate its average mark. Non passed exams count as 0.

```
select NumEt, avg (nvl (note, 0)) as Moy  
from R group by NumEt
```

/ nvl (x, y) : if x is not null then x else y*/*

R	Stuld	Moy
	12	36.67
	11	82.5

Générer des valeurs absentes

For each bar, give the drinkers frequenting it (including bars where nobody goes)

```
select Bar, Drinker from Frequents  
union
```

```
select Name, Null from Bars where Name not in (select Bar from Frequents)
```

The expected result is :

Bar	Drinker
Coogee Bay Hotel	Adam
Lord nelson	Gernot
Coogee Bay Hotel	John
Lord Nelson	John
Australia Hotel	John
Regent Hotel	Justin
Marble Bar	Justin
Royal Hotel	

This query is not correct : the schema `select Bar, Drinker from Frequents` is not compatible with `select Name, Null from Bars...` (Drinker belongs to Frequents keys, it cannot contain NULL values).

For each bar, give the drinkers frequenting it (including bars where nobody goes)

```
select B.Name, F.Drinker from Bars B left outer join  
        Frequents F on (B.name = F.bar)
```

For each tuple in Bars which do not satisfy the join relation with another tuple from Frequents, a tuple with NULL values is created.

What's the result of :

```
select F.Bar, F.Drinker  
from Bars B left outer join  
        Frequents F on (B.name = F.bar)
```

Attributes in the junction must be carefully chosen !

Another example :

Students (stuID, name, firstname)

Course (courseID, courseName)

Registrations (stuID, courseID)

For each student (ID, name and firstname), give the identifiers of courses she/he is registered (also give students who didn't register any course)

```
select S.stuID, S.name, S.lastname, R.courseID  
from Students S natural left outer join Registrations R
```

For each tuple in *Students* that do not satisfy the join condition (no tuple from registration), a tuple with NULL values for Registration is returned.

Another example :

For each student (stuID, name, firstname), give the identifiers and name of the courses she/he registered to (also give students not registered to any course)

```
select S.stuID, S.name, S.firstname, R.courseID, C.courseName  
from Student S natural left outer join Registrations R  
         natural left outer join Course C
```

The query below give the same result :

```
select S.stuID, S.name, S.firstname, R.courseID, C.courseName  
from Course C natural right outer join Registrations R  
         natural right outer join Students S
```

- Often, there are several ways of expressing the same query
 - Which one should be chosen ? Not an easy question
 - The queries are optimized by the DBMS, so developers do not need to worry about this...
 - However, with some DBMS, execution time can be critical